

Documentation: Pricing, Filtering, Weighting and Calibration for Deribit Inverse Options

(Module documentation for `calibration.py` and `inverse_fft_pricer.py`)

Abstract

This document describes, in mathematical detail, the pricing and calibration pipeline implemented in the modules `inverse_fft_pricer.py` and `calibration.py`. The pipeline ingests Deribit option snapshot data, filters for liquid quotes, builds weighted price residuals, and calibrates one of three models (Black, Heston, SVCJ) by solving a nonlinear least-squares problem where model prices are produced via the Carr–Madan FFT. Function names are referenced exactly as implemented.

Contents

1 Market objects and notation	2
1.1 Deribit snapshot observations	2
1.2 Per-expiry forward proxy and moneyness	2
2 Inverse option pricing via Carr–Madan FFT	2
2.1 Characteristic functions	2
2.2 Carr–Madan transform and damped call prices	3
2.3 FFT discretization (<code>carr_madan_call_fft</code>)	3
2.4 Inverse (coin) prices (<code>price_inverse_option</code>)	4
2.5 Caching of FFT grids	4
3 Dataset pricing for calibration (<code>price_dataframe</code>)	4
3.1 Dynamic log-strike centering (<code>_choose_fft_params_for_group</code>)	4
4 Filtering and cleaning (<code>filter_liquid_options</code>)	5
4.1 Required columns and numeric coercion	5
4.2 Liquidity and sanity constraints	5
4.3 Output	5
5 Weighting scheme (<code>WeightConfig</code> and <code>_weights</code>)	5
5.1 Weighted residual vector	5
5.2 Weight factorization	5
6 Calibration problem (<code>calibrate_model</code>)	6
6.1 Nonlinear least squares objective	6
6.2 Parameter constraints and unconstrained reparameterization	6
6.2.1 Black model	6
6.2.2 Heston model	7
6.2.3 SVCJ model	7
6.3 Initial guesses	7
6.4 Cache management	8
6.5 Fit diagnostics	8

7 End-to-end algorithm summary	8
8 Implementation cross-reference	8

1 Market objects and notation

1.1 Deribit snapshot observations

Fix a snapshot time t . Each row $i \in \{1, \dots, n\}$ corresponds to one listed European option instrument with:

- strike $K_i > 0$ in USD (column `strike`);
- time to maturity $T_i > 0$ in years (column `time_to_maturity`);
- option type $\tau_i \in \{\text{call}, \text{put}\}$ (column `option_type`);
- term futures mark F_i in USD per coin (column `futures_price`);
- bid/ask premium in coin units (columns `bid_price`, `ask_price`);
- liquidity proxies: vega V_i (column `vega`) and open interest OI_i (column `open_interest`).

The module `filter_liquid_options` constructs additional columns:

$$m_i := \frac{1}{2}(\text{bid}_i + \text{ask}_i) \quad (\text{mid_price_clean}), \quad (1)$$

$$s_i := \text{ask}_i - \text{bid}_i \quad (\text{spread}), \quad (2)$$

$$\text{rs}_i := \frac{s_i}{\max(m_i, \varepsilon)} \quad (\text{rel_spread}), \quad (3)$$

with a small numerical $\varepsilon > 0$.

1.2 Per-expiry forward proxy and moneyness

Options are grouped by the expiry key `expiry_datetime`. Within each expiry group g , `filter_liquid_options` defines a single forward proxy (“per-expiry median future”):

$$F_{0,g} := \text{median}\{F_i : i \in g\}, \quad (4)$$

stored as column `F0`. For each row $i \in g$ the module computes

$$\text{moneyness: } M_i := \frac{K_i}{F_{0,g}} \quad (\text{moneyness}), \quad \text{log-moneyness: } \kappa_i := \log\left(\frac{K_i}{F_{0,g}}\right) \quad (\text{log_moneyness}). \quad (5)$$

2 Inverse option pricing via Carr–Madan FFT

2.1 Characteristic functions

Let F_t denote the term futures price (USD per coin) and define

$$X_T := \log F_T. \quad (6)$$

For each model, the pricer assumes that the characteristic function

$$\phi_T(u) := \mathbb{E}[e^{iuX_T}], \quad u \in \mathbb{C}, \quad (7)$$

is available in closed (or semi-closed) form and is evaluable for complex arguments. In `inverse_fft_pricer.py` these are implemented as:

```

cf_black(u,T,F0,sigma),   cf_heston(u,T,F0,kappa,theta,sigma_v,rho,v0),
cf_svcj(u,T,F0,kappa,theta,sigma_v,rho,v0,lambda,sigma_y,ell_v,rho_j).

```

All three functions return $\phi_T(u)$ for vectorized complex u .

2.2 Carr–Madan transform and damped call prices

Let $C^{\text{USD}}(K)$ denote the (undiscounted) European call price in USD on the futures. Following Carr–Madan, define log-strike $k := \log K$ and the damped call transform

$$c_\alpha(k) := e^{\alpha k} C^{\text{USD}}(e^k), \quad \alpha > 0. \quad (8)$$

Under mild integrability conditions (existence of exponential moments), c_α admits the Fourier inversion

$$C^{\text{USD}}(K) = \frac{e^{-\alpha k}}{\pi} \int_0^\infty \Re(e^{-ivk} \Psi(v)) dv, \quad \Psi(v) := \frac{\phi_T(v - i(\alpha + 1))}{\alpha^2 + \alpha - v^2 + i(2\alpha + 1)v}. \quad (9)$$

This is implemented in `carr_madan_call_fft` by constructing $\Psi(v)$ on a frequency grid and recovering $C^{\text{USD}}(K)$ on a log-strike grid using an FFT.

2.3 FFT discretization (`carr_madan_call_fft`)

Choose an FFT configuration `FFTParams` with:

$$N \in 2^{\mathbb{N}}, \quad \eta > 0, \quad \alpha > 0, \quad b \in \mathbb{R}, \quad (10)$$

where N is the number of grid points, η the frequency spacing, α the damping factor and b the minimum log-strike on the output grid.

Frequency grid. Define

$$v_j := j\eta, \quad j = 0, 1, \dots, N - 1. \quad (11)$$

Log-strike grid. Define the dual spacing

$$\lambda := \frac{2\pi}{N\eta}, \quad (12)$$

and the log-strikes

$$k_u := b + u\lambda, \quad u = 0, 1, \dots, N - 1, \quad K_u := e^{k_u}. \quad (13)$$

Quadrature weights. Let w_j denote quadrature weights, either trapezoidal or Simpson, produced by `_quadrature_weights(N,use_simpson)`. The discretized integrand sequence is

$$x_j := e^{-iv_j b} \Psi(v_j) w_j \eta. \quad (14)$$

FFT and call grid. The FFT computes

$$y_u := \sum_{j=0}^{N-1} e^{-2\pi i j u / N} x_j, \quad (15)$$

and the damped call values are recovered via

$$c_\alpha(k_u) \approx \frac{1}{\pi} \Re(y_u). \quad (16)$$

Finally,

$$C^{\text{USD}}(K_u) \approx e^{-\alpha k_u} c_\alpha(k_u). \quad (17)$$

The function `carr_madan_call_fft` returns the arrays $(K_u)_{u=0}^{N-1}$ and $(C^{\text{USD}}(K_u))_{u=0}^{N-1}$.

2.4 Inverse (coin) prices (price_inverse_option)

Deribit quotes option premia in coin units. The pricer converts USD call prices to coin prices using:

$$C^{\text{coin}}(K, T; F_0, \theta) := \frac{C^{\text{USD}}(K, T; F_0, \theta)}{F_0}, \quad (18)$$

where F_0 is the current term futures price used as an exchange rate (USD per coin). In code this is the line `call_coin = call_prices_usd / F0`.

Puts are computed using inverse put–call parity as implemented:

$$C^{\text{coin}}(K, T) - P^{\text{coin}}(K, T) = 1 - \frac{K}{F_0} \implies P^{\text{coin}}(K, T) = C^{\text{coin}}(K, T) - \left(1 - \frac{K}{F_0}\right). \quad (19)$$

The public API entry point is

```
price_inverse_option(model, K, T, F0, params, option_type, fft_params).
```

2.5 Caching of FFT grids

To speed up calibration, the pricing grid $(K_u, C^{\text{USD}}(K_u))$ is cached by an LRU cache:

$$\text{_cached_pricing_grid}: (m, T, F_0, \theta, \text{FFTParams}) \mapsto (K_u, C^{\text{USD}}(K_u)). \quad (20)$$

Thus repeated evaluations at the same maturity T , forward F_0 , model name m , parameters θ and FFT grid reuse the FFT output rather than recomputing it.

3 Dataset pricing for calibration (price_dataframe)

Given a filtered DataFrame \mathcal{D} , `price_dataframe` produces the vector of model prices $\mathbf{P}^{\text{model}} \in \mathbb{R}^n$ in coin units by:

1. grouping rows by `expiry_datetime` and `option_type`;
2. for each group g , extracting a representative maturity and forward:

$$T_g := \text{median}\{T_i : i \in g\}, \quad F_{0,g} := \text{median}\{F_{0,i} : i \in g\}; \quad (21)$$

3. calling `price_inverse_option` once per group with the vector of strikes $\{K_i\}_{i \in g}$.

3.1 Dynamic log-strike centering (_choose_fft_params_for_group)

FFT accuracy requires that the strike grid $\{K_u\}$ covers the strikes of interest. The helper `_choose_fft_params_for_group` keeps (N, η, α) fixed, computes

$$\lambda = \frac{2\pi}{N\eta}, \quad k_{\text{center}} := \log(\text{median}\{K_i : i \in g\}), \quad (22)$$

and selects the left endpoint

$$b_g := k_{\text{center}} - \frac{N\lambda}{2}. \quad (23)$$

Then the grid $k_u = b_g + u\lambda$ is centered on the group's median strike. This reduces interpolation error and helps ensure that the grid spans the relevant strikes.

4 Filtering and cleaning (`filter_liquid_options`)

4.1 Required columns and numeric coercion

The function `filter_liquid_options` requires the columns

```
{currency, option_type, strike, time_to_maturity, bid_price, ask_price,
futures_price, vega, open_interest, expiry_datetime}.
```

Columns used in computations are coerced to numeric (invalid parses become NaN) before filtering.

4.2 Liquidity and sanity constraints

Let i index a row. The filtration applies the following constraints:

$$(\text{TTM domain}) \quad T_i > 0, \quad T_i \geq T_{\min}, \quad (T_i \leq T_{\max} \text{ if provided}). \quad (24)$$

$$(\text{Bid/ask present}) \quad \text{bid}_i > 0, \quad \text{ask}_i > 0, \quad \text{ask}_i \geq \text{bid}_i. \quad (25)$$

$$(\text{Vega / OI}) \quad V_i \geq V_{\min}, \quad \text{OI}_i \geq \text{OI}_{\min}. \quad (26)$$

$$(\text{Relative spread}) \quad \text{rs}_i = \frac{\text{ask}_i - \text{bid}_i}{\max(m_i, \varepsilon)} \leq \text{rs}_{\max} \quad \text{if provided.} \quad (27)$$

$$(\text{Moneyness window}) \quad M_i = \frac{K_i}{F_{0,g}} \in [M_{\min}, M_{\max}] \quad \text{if provided.} \quad (28)$$

$$(\text{Option type}) \quad \tau_i \in \{\text{call}, \text{put}\} \text{ after lowercasing.} \quad (29)$$

Optionally, synthetic underlyings with names beginning `SYN`. may be dropped via `drop_synthetic_underlyings=True`.

4.3 Output

The returned DataFrame is sorted by `expiry_datetime`, `strike` and `option_type` and includes the derived columns $(m_i, s_i, \text{rs}_i, F_{0,i}, M_i, \kappa_i)$ described above.

5 Weighting scheme (`WeightConfig` and `_weights`)

5.1 Weighted residual vector

Calibration is performed by nonlinear least squares on a residual vector whose components are

$$r_i(\theta) := w_i \left(P_i^{\text{model}}(\theta) - P_i^{\text{mkt}} \right), \quad (30)$$

where:

- $P_i^{\text{mkt}} := m_i$ is the cleaned mid price (`mid_price_clean`);
- $P_i^{\text{model}}(\theta)$ is produced by `price_dataframe` using `price_inverse_option`.

5.2 Weight factorization

The weight w_i is constructed multiplicatively from three components (spread, vega, open interest) as implemented in `_weights`:

$$w_i = \left(s_i + \varepsilon_s \right)^{-p_s \cdot \mathbf{1}_{\text{spread}}} \left(V_i + \varepsilon_o \right)^{p_v \cdot \mathbf{1}_{\text{vega}}} \left(\text{OI}_i + \varepsilon_o \right)^{p_o \cdot \mathbf{1}_{\text{oi}}}, \quad (31)$$

where:

- $s_i = \text{ask}_i - \text{bid}_i$ is the bid–ask spread;
- V_i is the snapshot vega;
- OI_i is the snapshot open interest;
- $\varepsilon_s = \text{eps_spread}$ and $\varepsilon_o = \text{eps_other}$ are small stabilizers;
- exponents are configured by `spread_power` (p_s), `vega_power` (p_v), and `oi_power` (p_o);
- the indicator flags $\mathbf{1}_{\text{spread}}, \mathbf{1}_{\text{vega}}, \mathbf{1}_{\text{oi}} \in \{0, 1\}$ switch components on/off via `use_spread`, `use_vega`, `use_open_interest`.

Finally, weights may be capped to prevent dominance:

$$w_i \leftarrow \min\{w_i, w_{\max}\} \quad \text{with } w_{\max} = \text{cap}. \quad (32)$$

If w_i is non-finite it is set to 0 (effectively dropping the observation).

Interpretation. With $p_s = 1$, r_i behaves like a (scaled) error normalized by the spread. With $p_v = p_o = \frac{1}{2}$, weights scale as $\sqrt{\text{vega} \times \text{OI}}$.

6 Calibration problem (`calibrate_model`)

6.1 Nonlinear least squares objective

Given a filtered dataset \mathcal{D} , define the residual map

$$\mathbf{r}(\theta) := (r_1(\theta), \dots, r_n(\theta))^{\top} \in \mathbb{R}^n, \quad (33)$$

with r_i as in eq. (30). The calibration estimates parameters by

$$\hat{\theta} \in \arg \min_{\theta \in \Theta} \frac{1}{2} \|\mathbf{r}(\theta)\|_2^2. \quad (34)$$

In code, `scipy.optimize.least_squares` (method "trf") solves eq. (34).

6.2 Parameter constraints and unconstrained reparameterization

The optimizer operates on an unconstrained vector $x \in \mathbb{R}^d$. Each model uses a bijection (or a near-bijection) $\theta = T(x)$ to enforce constraints.

6.2.1 Black model

Parameters: $\theta = (\sigma)$ with $\sigma > 0$.

$$x = \log \sigma \iff \sigma = e^x.$$

Implemented by `_pack_black` and `_unpack_black`.

6.2.2 Heston model

Parameters: $\theta = (\kappa, \theta, \sigma_v, \rho, v_0)$ with

$$\kappa > 0, \quad \theta > 0, \quad \sigma_v > 0, \quad v_0 > 0, \quad \rho \in (-1, 1).$$

Transform:

$$x_1 = \log \kappa, \quad \kappa = e^{x_1}, \quad (35)$$

$$x_2 = \log \theta, \quad \theta = e^{x_2}, \quad (36)$$

$$x_3 = \log \sigma_v, \quad \sigma_v = e^{x_3}, \quad (37)$$

$$x_4 = \operatorname{arctanh}(\rho), \quad \rho = \tanh(x_4), \quad (38)$$

$$x_5 = \log v_0, \quad v_0 = e^{x_5}. \quad (39)$$

Implemented by `_pack_heston` and `_unpack_heston`.

6.2.3 SVCJ model

Parameters extend Heston by jump parameters $\lambda, \ell_y, \sigma_y, \ell_v, \rho_j$:

$$\lambda > 0, \quad \sigma_y > 0, \quad \ell_v > 0,$$

while $\ell_y \in \mathbb{R}$ is unconstrained and ρ_j is constrained indirectly to preserve finiteness of the jump MGF.

Key stability constraint. The SVCJ characteristic function implementation requires

$$1 - \ell_v \rho_j > 0, \quad (40)$$

to keep the denominator of the jump MGF positive (and thus finite) at the relevant complex arguments.

Reparameterization used in code. The module introduces an auxiliary bounded scalar

$$\psi := \ell_v \rho_j \in (-0.99, 0.99), \quad (41)$$

and optimizes over x_{10} such that

$$\psi = 0.99 \tanh(x_{10}), \quad \rho_j = \frac{\psi}{\ell_v}. \quad (42)$$

This ensures $\psi < 1$ and, in practice, supports eq. (40) for most parameter values. Other parameters are transformed as:

$$\lambda = e^{x_6}, \quad \ell_y = x_7, \quad \sigma_y = e^{x_8}, \quad \ell_v = e^{x_9}. \quad (43)$$

Together with the Heston transforms for $(\kappa, \theta, \sigma_v, \rho, v_0)$, this defines the SVCJ mapping, implemented by `_pack_svcj` and `_unpack_svcj`.

6.3 Initial guesses

If the dataset contains `implied_volatility`, the initializer sets an “ATM” proxy

$$\sigma_{\text{atm}} := \text{median}\{\text{implied_volatility}_i\}, \quad (44)$$

and uses $\theta_{\text{init}} = \sigma_{\text{atm}}^2$ for variance-level parameters. Otherwise a fallback $\sigma_{\text{atm}} = 0.6$ is used. Implemented in `_default_initial_params`.

6.4 Cache management

Before running the optimizer, `calibrate_model` may clear the FFT grid cache by calling `clear_fft_cache()`, which in turn invokes `_cached_pricing_grid.cache_clear()`.

6.5 Fit diagnostics

After obtaining $\hat{\theta}$, the code computes model prices on the fitting set and reports:

$$\text{RMSE} := \sqrt{\frac{1}{n} \sum_{i=1}^n (P_i^{\text{model}}(\hat{\theta}) - P_i^{\text{mkt}})^2}, \quad \text{MAE} := \frac{1}{n} \sum_{i=1}^n |P_i^{\text{model}}(\hat{\theta}) - P_i^{\text{mkt}}|. \quad (45)$$

7 End-to-end algorithm summary

Algorithm 1 Calibration pipeline (implemented by `filter_liquid_options` and `calibrate_model`)

Require: Raw snapshot DataFrame \mathcal{D}_{raw} , model name $m \in \{\text{black}, \text{heston}, \text{svcj}\}$

- 1: $\mathcal{D} \leftarrow \text{filter_liquid_options}(\mathcal{D}_{\text{raw}}, \text{liquidity thresholds})$
 - 2: Construct market mid prices \mathbf{P}^{mkt} and weights \mathbf{w} via `_weights` and `WeightConfig`
 - 3: Choose initial parameters $\theta^{(0)}$ via `_default_initial_params`
 - 4: Map to unconstrained $x^{(0)} = \text{_pack_*}(\theta^{(0)})$
 - 5: Define residual function $x \mapsto \mathbf{r}(T(x))$ where model prices use `price_dataframe` o `price_inverse_option`
 - 6: Solve $\hat{x} \in \arg \min_x \frac{1}{2} \|\mathbf{r}(T(x))\|_2^2$ using `least_squares(method="trf")`
 - 7: Return $\hat{\theta} = T(\hat{x})$ and diagnostics (`CalibrationResult`)
-

8 Implementation cross-reference

- **Filtering:** `filter_liquid_options` (adds `mid_price_clean`, `spread`, `rel_spread`, `F0`, `moneyness`, `log_moneyness`).
- **Weights:** `WeightConfig` and `_weights` implement eq. (31).
- **Pricing:** `price_dataframe` groups by expiry and calls `price_inverse_option`.
- **FFT grid selection:** `_choose_fft_params_for_group` selects b_g for centered strike coverage.
- **Carr–Madan engine:** `carr_madan_call_fft` implements eq. (9).
- **Inverse conversion:** `price_inverse_option` implements eq. (18) and eq. (19).
- **Calibration:** `calibrate_model` solves eq. (34) with parameter transforms `_pack_*` / `_unpack_*`.